

# Parallel Algorithms

Jongseo Lee (KAIST)

2023-01-10

# Note

## Goals & Non-Goals

- Goals
  - Explore about two recent models of parallelism
- Non-Goals
  - Understand everything in the textbook

# Massively Parallel Computing Model (MPC)

## Definition

- **Massively Parallel Computing Model (MPC)**
  - What does the MPC consist of?
  - How does the input given to the MPC?
  - How does the MPC output the answer?
  - How does the MPC perform computation?
  - Which parameters of the MPC are important?

# Massively Parallel Computing Model (MPC)

## Definition

- The **Massively Parallel Computing Model (MPC)** consists of:
  - The input data size  $N$
  - The number of machines  $M$  (machines will communicate with each other)
  - The memory size,  $s$  words, each machine can hold
- We are mostly interested if  $M \cdot s = \Theta(N)$ ,  $M \cdot s = O(N \text{ polylog } N)$ , etc.
  - Note that  $M \cdot s$  is the amount of the available resources

# Massively Parallel Computing Model (MPC)

## Definition

- Input
  - The input data is split across the  $M$  machines arbitrarily
- Output
  - Some machines will be the **output machines**.
  - The answer will be in the output machines.

# Massively Parallel Computing Model (MPC)

## Definition

Example - An integer array  $A[]$  is given in the form of  $(i, a_i)$  across the machines

$M_0$	(0, 2)	(2, 2)	(3, 3)	(8, 1)	(10, 3)	
-------	--------	--------	--------	--------	---------	--

$M_1$	(4, 0)	(6, 1)	(1, 0)			
-------	--------	--------	--------	--	--	--

$M_2$	(7, 0)	(11, 4)	(5, 1)			
-------	--------	---------	--------	--	--	--

# Massively Parallel Computing Model (MPC)

## Definition

Example - The output data, prefix sums of the given array are stored across the machines

M <sub>0</sub>	(0, 2)	(1, 2)	(2, 4)	(3, 7)		
----------------	--------	--------	--------	--------	--	--

M <sub>1</sub>	(4, 7)	(5, 8)	(6, 9)	(7, 9)		
----------------	--------	--------	--------	--------	--	--

M <sub>2</sub>	(8, 10)	(9, 12)	(10, 15)	(11, 19)		
----------------	---------	---------	----------	----------	--	--

# Massively Parallel Computing Model (MPC)

## Definition

- In each round of the computation,
  1. each machine performs some computation on the local data
  2. each machine sends/receives messages to any other machine
    - Note: A machine can send/receive at most  $s$  words per round



# Massively Parallel Computing Model (MPC)

## Definition

- Main Parameters
  - The number of machines,  $M$  (mostly omitted)
  - The number of communication rounds an algorithm
  - The size of local memory  $s$ 
    - Problems are easier to solve with larger  $s$ . What if  $s \gg N$ ?
    - Typically,  $s = O(N^\varepsilon)$  for  $\varepsilon < 1$ .

# Massively Parallel Computing Model (MPC)

## Definition

### Definition.

An MPC algorithm is **sublinear** if  $s \ll N$ . Formally,  $s = N^\varepsilon$  for some constant  $\varepsilon < 1$ .

### Example.

An MPC algorithm with  $s = n^{1+\delta}$  for some  $\delta < 1$  on dense graphs is considered to be sublinear.

# Massively Parallel Computing Model (MPC)

Example: Computing Prefix Sum

- Input:  $n$  pair of integers  $(0, a_0), (1, a_1), \dots, (N - 1, a_{N-1})$
- Output:  $(i, \sum_{j=0}^i a_j)$  for  $0 \leq i < N$ .
- In this example, we assume  $M, s = \Theta(N^{0.5})$ .

# Massively Parallel Computing Model (MPC)

Example: Computing Prefix Sum

- Input:  $n$  pair of integers  $(0, a_0), (1, a_1), \dots, (N-1, a_{N-1})$
- Output:  $(i, \sum_{j=0}^i a_j)$  for  $0 \leq i < N$ .

2	0	2	3	0	1	1	0	1	2	3	4
2	2	4	7	7	8	9	9	10	12	15	19

# Massively Parallel Computing Model (MPC)

Example: Computing Prefix Sum

Initial State

$M_0$	(0, 2)	(2, 2)	(3, 3)	(8, 1)	(10, 3)	
-------	--------	--------	--------	--------	---------	--

$M_1$	(4, 0)	(6, 1)	(1, 0)			
-------	--------	--------	--------	--	--	--

$M_2$	(7, 0)	(11, 4)	(5, 1)			
-------	--------	---------	--------	--	--	--

# Massively Parallel Computing Model (MPC)

Example: Computing Prefix Sum

Communication Round 1: Send  $(i, A_i)$  to machine  $\left\lfloor \frac{Mi}{N} \right\rfloor$

M <sub>0</sub>	(0, 2)	(1, 0)	(2, 2)	(3, 3)		
----------------	--------	--------	--------	--------	--	--

M <sub>1</sub>	(4, 0)	(5, 1)	(7, 0)	(6, 1)		
----------------	--------	--------	--------	--------	--	--

M <sub>2</sub>	(9, 2)	(11, 4)	(8, 1)	(10, 3)		
----------------	--------	---------	--------	---------	--	--

# Massively Parallel Computing Model (MPC)

## Example: Computing Prefix Sum

Computation Round 1: **Sort** and compute prefix sum in each machine

M <sub>0</sub>	(0, 2)	(1, 0)	(2, 2)	(3, 3)		
----------------	--------	--------	--------	--------	--	--

M <sub>1</sub>	(4, 0)	(5, 1)	(6, 1)	(7, 0)		
----------------	--------	--------	--------	--------	--	--

M <sub>2</sub>	(8, 1)	(9, 2)	(10, 3)	(11, 4)		
----------------	--------	--------	---------	---------	--	--

# Massively Parallel Computing Model (MPC)

## Example: Computing Prefix Sum

Computation Round 1: Sort and **compute prefix sum** in each machine

M <sub>0</sub>	(0, 2)	(1, 2)	(2, 4)	(3, 7)		
----------------	--------	--------	--------	--------	--	--

M <sub>1</sub>	(4, 0)	(5, 1)	(6, 2)	(7, 2)		
----------------	--------	--------	--------	--------	--	--


M <sub>2</sub>	(8, 1)	(9, 3)	(10, 6)	(11, 10)		
----------------	--------	--------	---------	----------	--	--



# Massively Parallel Computing Model (MPC)

## Example: Computing Prefix Sum

Communication Round 2: Send total sum of machine  $i$  to machine  $j(j > i)$



# Massively Parallel Computing Model (MPC)

## Example: Computing Prefix Sum

Computation Round 2: Add the numbers received into prefix sums

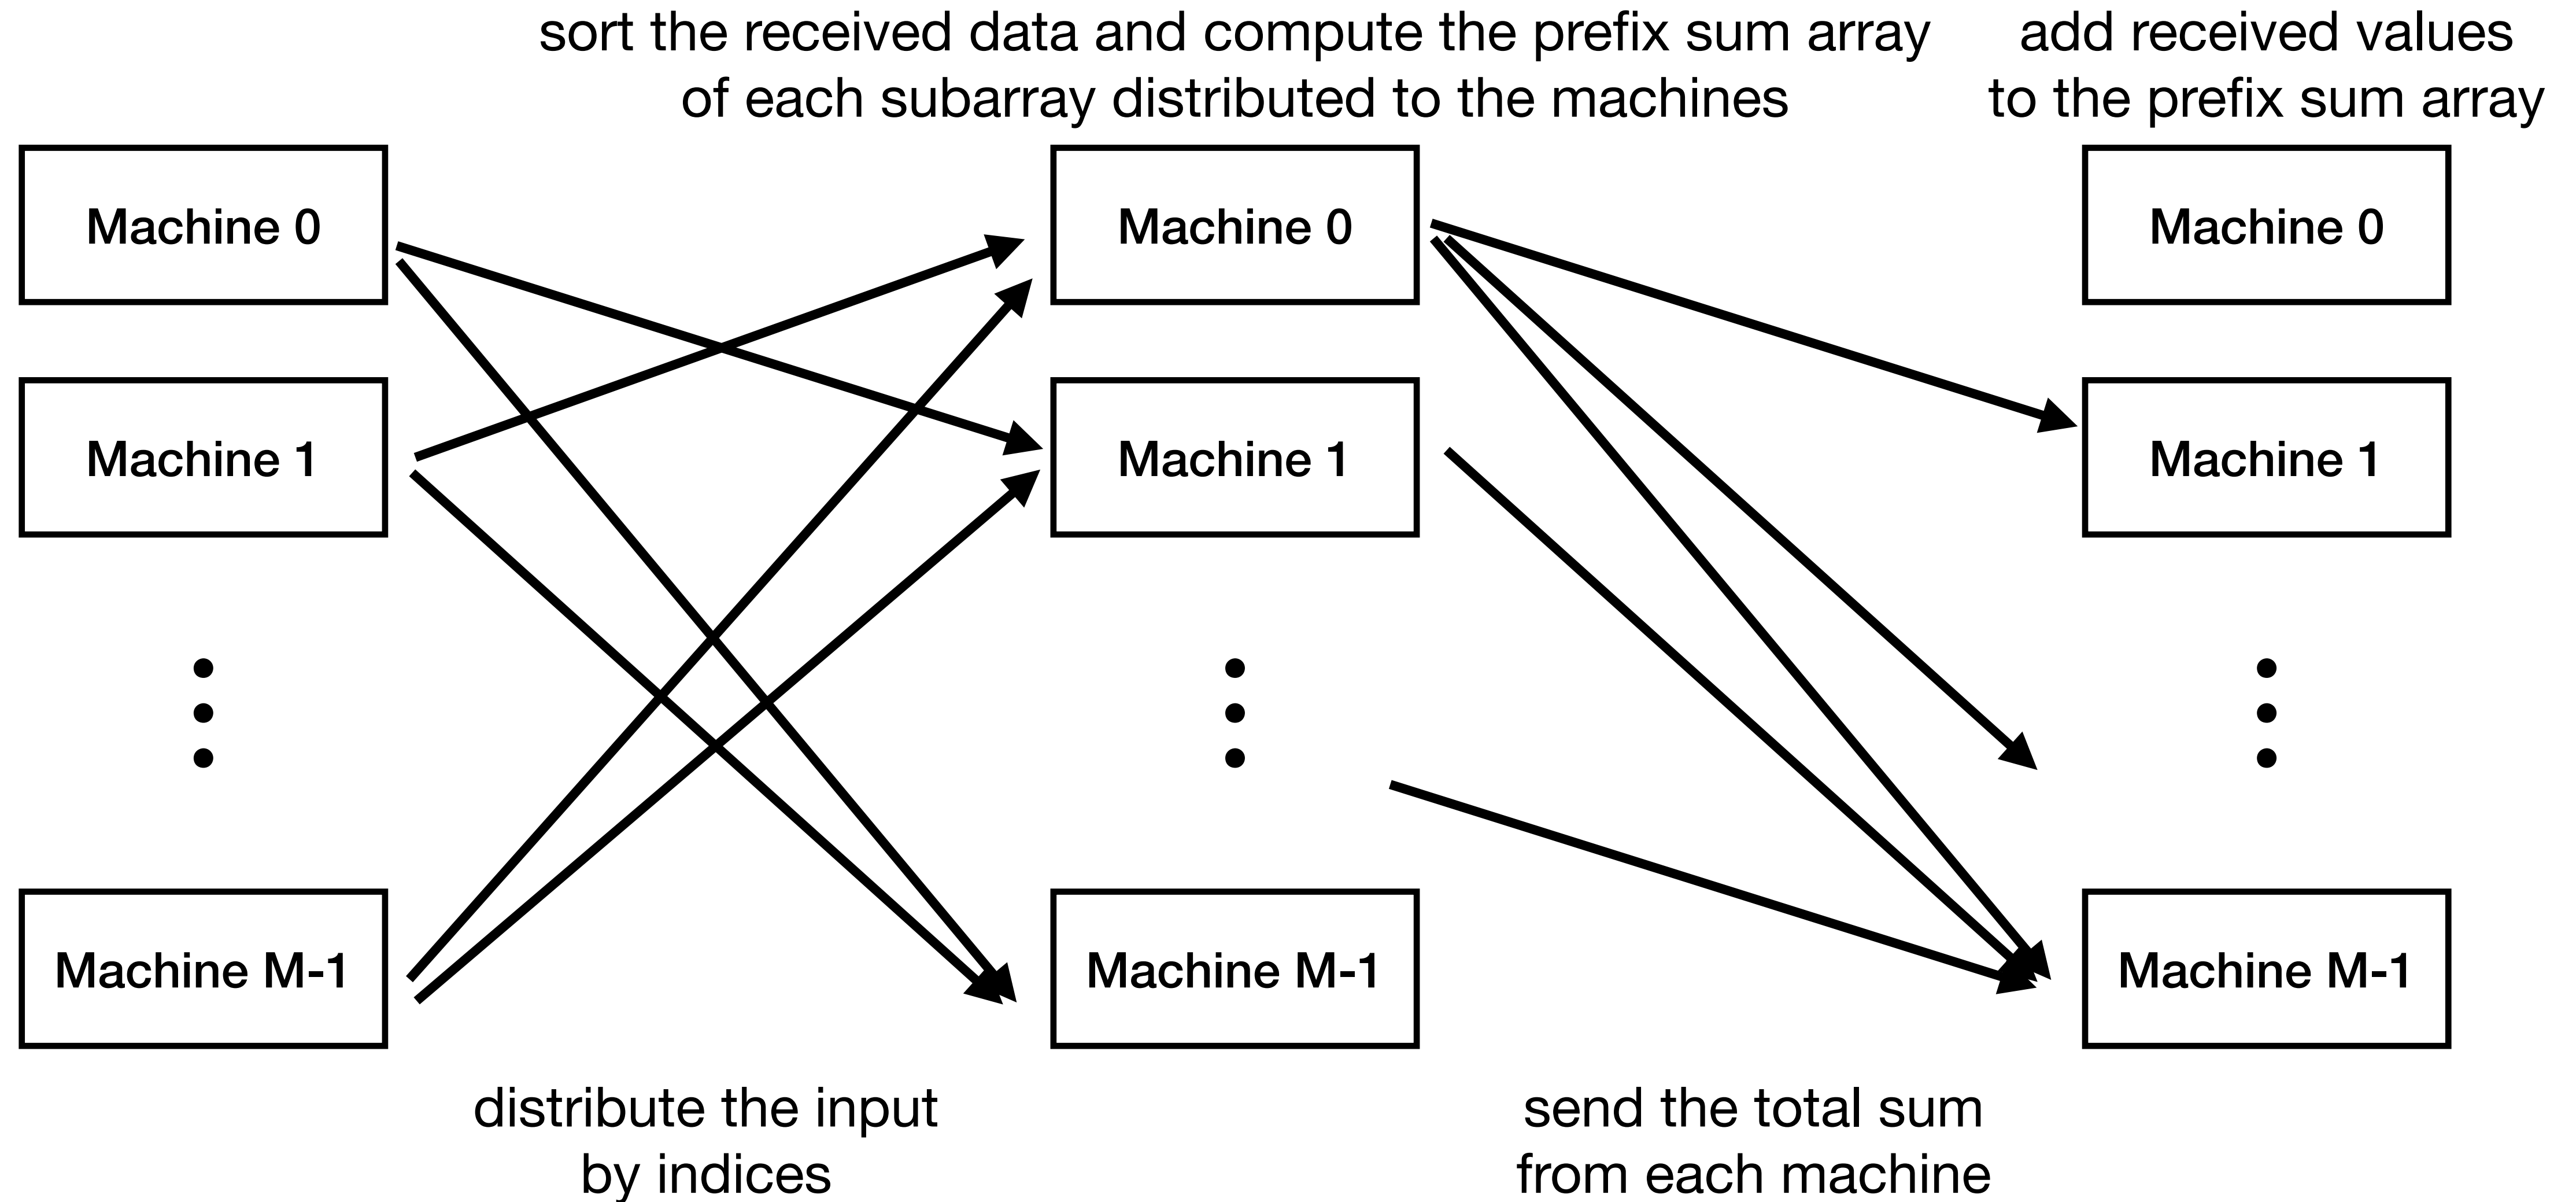
$M_0$	(0, 2)	(1, 2)	(2, 4)	(3, 7)		
-------	--------	--------	--------	--------	--	--

$M_1$	(4, 7)	(5, 8)	(6, 9)	(7, 9)		
-------	--------	--------	--------	--------	--	--

$M_2$	(8, 10)	(9, 12)	(10, 15)	(11, 19)		
-------	---------	---------	----------	----------	--	--

# Massively Parallel Computing Model (MPC)

## Example: Computing Prefix Sum



# Massively Parallel Computing Model (MPC)

Example: Computing Prefix Sum

- By taking  $M = N^{0.5}$ ,  $s = 2 \times N^{0.5}$ , we used
  - $O(N^{0.5})$  machines each of them having  $O(N^{0.5})$  words of memory
  - $O(1)$  rounds of communications, with total  $O(N)$  words

# Massively Parallel Computing Model (MPC)

## Algorithms - Maximal Matching

### Maximal Matching

- Instance

A graph  $G = (V, E)$

- Question

Return a matching  $M \subseteq E$  with every edge in  $E \setminus M$  having an endpoint in  $M$

# Massively Parallel Computing Model (MPC)

## Algorithms - Maximal Matching

### Overview

- Repeat until the number of remaining edges fit into a single machine:
  - Put a subset of edges into a single machine
  - Compute a maximal matching
  - remove matched vertices

# Massively Parallel Computing Model (MPC)

## Algorithms - Maximal Matching

The algorithm (Note:  $s = n^{1+\varepsilon}$ )

- Let  $G_0 := G$
- For round  $r = 0, 1, \dots, R - 1$  (Note:  $E(G_r) = \emptyset$ )
  - Each machine marks each of its edges independently with probability  $p = \frac{n^{1+\varepsilon}}{2m}$
  - Each machine sends marked edges to machine 0
  - Machine 0 computes a maximal matching  $M_r$  and announce matched vertices
  - Remove edges having a matched endpoint and set  $G_{r+1}$  as a current graph

# Massively Parallel Computing Model (MPC)

## Algorithms - Maximal Matching

Fix a round and let  $m$  be the number of edges at the start of the round.

### Lemma.

With high probability, the number of marked edges fit into a single machine

### Lemma.

With high probability, the number of remaining edges is at most  $\frac{10m}{n^\epsilon}$



# Massively Parallel Computing Model (MPC)

## Algorithms - Maximal Matching

### Theorem.

The algorithm terminates in  $R = O(1/\varepsilon)$  rounds

### Proof.

There are at most  $n^2$  edges initially

After  $R \leq \log_{n^\varepsilon} n^2 = O(1/\varepsilon)$  rounds, due to the second Lemma, there are at most  $s = n^{1+\varepsilon}$  edges left.

# Massively Parallel Computing Model (MPC)

## Algorithms - Connected Components

### Connected Components (ConnComp)

- Instance

An undirected graph  $G = (V, E)$

- Question

Which vertices are in the same connected component?

A solution is a labeling of vertices s.t.  $\ell(u) = \ell(v)$  iff there's a path from  $u$  to  $v$

# Massively Parallel Computing Model (MPC)

## Algorithms - Connected Components

Idea for handling sparse graphs: Graph Exponentiation


- Operation: Connect every vertex with vertices of distance  $\leq 2$
- Repeating the operation  $O(\log D)$  steps, each connected component becomes a clique
  - $D$ : diameter of  $G$ , the length of the longest shortest path
- Problem: required memory of a single machine can be up to  $\Omega(n^2)$

# Massively Parallel Computing Model (MPC)

## Algorithms - Connected Components

Idea for handling dense graphs: Label Contraction

- Operation
  - Mark each vertex independently with prob.  $p$
  - Relabel each unmarked vertex with a marked neighbor
  - The denser the graph, the smaller the value of  $p$  required for the contraction



# Massively Parallel Computing Model (MPC)

## Algorithms - Connected Components

The Algorithm: Putting them together (Rough Sketch)

- Repeat the following phase:
  1. Perform the graph exponentiation, collapse the vertices into a supernode
  2. Perform label contraction
- Choosing the appropriate  $p$  allows us to obtain a sublinear algorithm having  $O(\log D \cdot \log \log_{m/n} n)$  round complexity
  - Each phase takes  $O(\log D)$  rounds
  - $O(\log \log_{m/n} n)$  phases suffice

# Massively Parallel Computing Model (MPC)

What the MPC can't do

## 2-Cycle Problem

- Instance: A graph that is either one cycle or the union of two cycles
- Question: Determine if the graph is one cycle or the union of two cycles

## 2-Cycle Conjecture

Any sublinear MPC algorithm for the 2-Cycle problem requires  $\Omega(\log n)$  rounds

Furthermore, it is conjectured to hold even for the problem of distinguishing

- A cycle of length  $n$
- Or the two cycles of length  $n/2$

# Massively Parallel Computing Model (MPC)

What the MPC can't do

**Theorem.** Assuming 2-Cycle conjecture, any sublinear MPC for undirected connectivity problem requires  $\Omega(\log n)$  rounds

Proof. Reduction from 2-Cycle problem to undirected connectivity problem

- consisting of one cycle  $\rightarrow$  connected
- consisting of two cycles  $\rightarrow$  disconnected

# Massively Parallel Computing Model (MPC)

What the MPC can't do

**Component-stable MPC algorithms** for a graph problem

- The output of each vertex depends only on the connected component that it belongs to



# Massively Parallel Computing Model (MPC)

What the MPC can't do

**Theorem.** Assuming 2-Cycle conjecture, following holds for sublinear component-stable MPC algorithms

- Maximal matching requires  $\Omega(\log \log n)$  rounds (including its constant approximation)
- For every constant  $c$ ,  $c$ -coloring of a cycle requires  $\Omega(\log \log^* n)$  rounds
- Any constant approximation for Vertex Cover requires  $\Omega(\log \log n)$  rounds
- (Using more assumptions)  $(\Delta + 1)$ -coloring of a graph requires  $\Omega(\sqrt{\log \log n})$  rounds

# Adaptive Massively Parallel Model (AMPC)

## Definition

The **Adaptive Massively Parallel Model (AMPC)** is basically the MPC model but with shared random access memory

- In the  $i$ -th round, each machine can read data from random access memory  $D_{i-1}$  and write to  $D_i$ 
  - $D_i$  s are common for all machines
- Each machine can perform  $\leq S$  reads and  $\leq S$  writes in a single round
  - Note:  $S$  is different from  $s$

Clearly, the AMPC is at least as strong as the MPC.

# Adaptive Massively Parallel Model (AMPC)

## Example

- Suppose there's a function  $g : X \rightarrow X$  and  $\{(x, g(x)) : x \in X\} \subseteq D_{i-1}$ .
- Then, given  $y \in X$  and  $k = O(S)$ , a machine can compute  $g^k(y)$ .

# Adaptive Massively Parallel Model (AMPC)

What the AMPC can do better than the MPC

For now, **2-Cycle problem** refers to the problem of distinguishing between:

- A single cycle on  $n$  vertices
- Two cycles on  $n/2$  vertices each

We aim to derive an AMPC algorithm solving 2-Cycle problem in  $O(1)$  rounds with high probability

# Adaptive Massively Parallel Model (AMPC)

What the AMPC can do better than the MPC

The Algorithm:  $s = O(n^\epsilon)$  (sublinear) and takes  $O(1/\epsilon) = O(1)$  rounds

## ALGORITHM 1: SHRINK( $G = (V, E), \delta, t$ )

- (1) For  $i := 1, \dots, t$ :
  - (a) Sample each vertex independently with probability  $n^{-\delta/2}$ . Denote the set of sampled vertices by  $M$ . Randomly distribute the vertices of  $M$  to the machines.
  - (b) For each sampled vertex  $v$  traverse the cycle in each direction until a sampled vertex is reached. Let  $l_v$  and  $r_v$  be the sampled vertices, where the each of two traversals finishes.
- (2) Return the graph  $(M, \{xl_x \mid x \in M\} \cup \{xr_x \mid x \in M\})$ .

## ALGORITHM 2: 2-CYCLE( $G = (V, E)$ )

- (1)  $G' := \text{SHRINK}(G, \epsilon, O(1/\epsilon))$
- (2) Solve the 2-CYCLE problem on  $G'$ , which has size  $O(n^\epsilon)$  w.h.p. on a single machine.

# Adaptive Massively Parallel Model (AMPC)

What the AMPC can do better than the MPC

## Lemma.

Let  $G$  be a graph consisting of cycles and  $n := |V(G)|$ . Fix an iteration from  $SHRINK(G, \varepsilon, O(1/\varepsilon))$ .

If the size of a cycle was  $\Omega(n^\varepsilon)$  at the beginning of the iteration, its size shrinks by at least  $\Omega(n^{\varepsilon/2})$  times after the iteration w.h.p.

## Lemma.

After the shrink operation, the length of each cycle becomes  $O(n^\varepsilon)$  w.h.p.

## Lemma.

Total communication of each machine is  $O(n^\varepsilon)$  in each round w.h.p.

# Adaptive Massively Parallel Model (AMPC)

What the AMPC can do better than the MPC

Problem	AMPC	MPC
Connectivity	$O(\log \log_{m/n} n)$	$O(\log D + \log \log_{m/n} n)$ [10]
Minimum spanning tree	$O(\log \log_{m/n} n)$	$O(\log n)$
2-edge connectivity	$O(\log \log_{m/n} n)$	$O(\log D \cdot \log \log_{m/n} n)$ [2]
Maximal independent set	$O(1)$	$\tilde{O}(\sqrt{\log n})$ [25]
2-CYCLE	$O(1)$	$O(\log n)$
Forest Connectivity	$O(1)$	$O(\log D \cdot \log \log_{m/n} n)$ [2]

Fig. 1. Round complexities of our algorithms in the AMPC model compared to the state-of-the-art MPC algorithms.  $D$  denotes the diameter of the input graph. We consider the setting where the space per machine is sublinear in the number of vertices, that is  $S = O(n^\epsilon)$  for constant  $\epsilon < 1$ .

# Adaptive Massively Parallel Model (AMPC)

What the AMPC can't do

**Theorem.** (Note: 2-Cycle problem refers to the original 2-Cycle problem)

- Any deterministic AMPC algorithm for 2-Cycle problem takes  $\geq \frac{1}{2} \log_S(n/2)$  rounds
  - If  $S = n^\epsilon$ , the number of rounds is  $\Omega(1/\epsilon)$
- Any *good (informal)* randomized AMPC algorithm for 2-Cycle takes  $\geq \frac{1}{2} \log_S(n/2)$  rounds
  - If  $S = n^\epsilon$ , the number of rounds is  $\Omega(1/\epsilon)$



# References

- Computational Intractability: A Guide to Algorithmic Lower Bounds  
<https://hardness.mit.edu/>
- S. Behnezhad et al. Massively Parallel Computation via Remote Memory Access. *ACM Transactions on Parallel Computing*, Vol. 8, No. 3
- Massively Parallel Algorithms  
<http://people.csail.mit.edu/ghaffari/MPA19/Notes/MPA.pdf>